

```

# =====
# MNIST Image Processing & Analysis Assignment
# Framework: Google Colab
# Dataset: MNIST Handwritten Digits
# Algorithms: KMeans Clustering (unsupervised), Logistic Regression (supervised)
# =====

# -----
# 1 Install / Import Libraries
# -----
# Google Colab usually has these pre-installed
import numpy as np
import matplotlib.pyplot as plt
from sklearn.cluster import KMeans
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score, confusion_matrix
from sklearn.preprocessing import StandardScaler
import seaborn as sns
import tensorflow as tf

# -----
# 2 Load MNIST Dataset
# -----
# MNIST comes preloaded with TensorFlow
(x_train, y_train), (x_test, y_test) =
tf.keras.datasets.mnist.load_data()


print("Training set shape:", x_train.shape)
print("Test set shape:", x_test.shape)

Downloading data from https://storage.googleapis.com/tensorflow/tf-
keras-datasets/mnist.npz
11490434/11490434 ————— 0s 0us/step
Training set shape: (60000, 28, 28)
Test set shape: (10000, 28, 28)


# -----
# 3 Show Example of Raw Data
# -----
# Display first 3 raw images
for i in range(3):
    plt.imshow(x_train[i], cmap='gray')
    plt.title(f"Raw Image - Label: {y_train[i]}")
    plt.show()

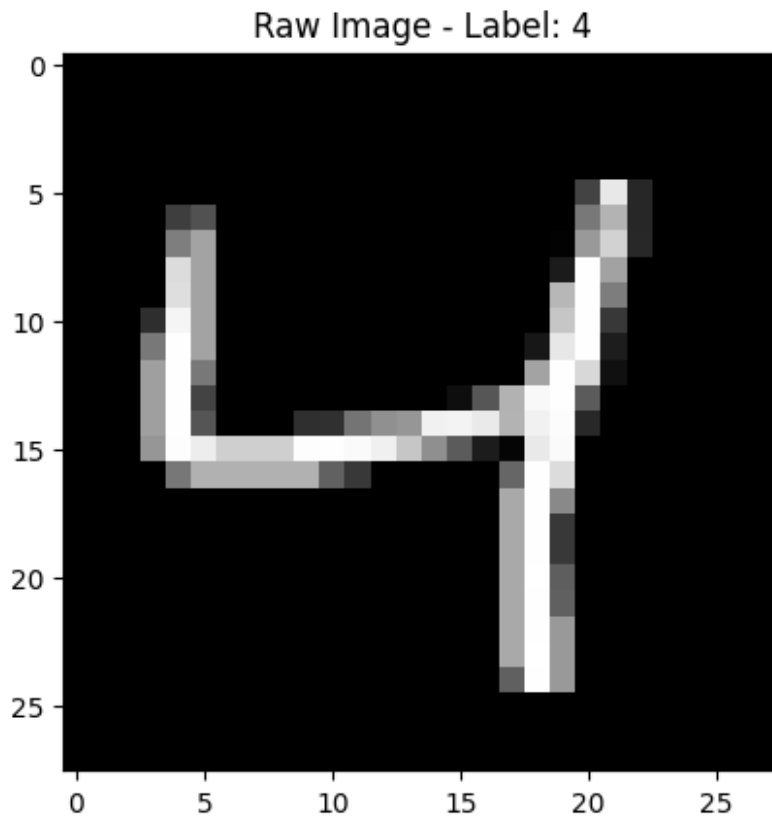
```

Raw Image - Label: 5



Raw Image - Label: 0





```
# -----
# 4 Preprocessing (Handling Unstructured Data)
# -----
# Normalize pixel values to 0-1
x_train_norm = x_train / 255.0
x_test_norm = x_test / 255.0

# Flatten images for ML algorithms (28x28 -> 784 features)
x_train_flat = x_train_norm.reshape(len(x_train_norm), -1)
x_test_flat = x_test_norm.reshape(len(x_test_norm), -1)

print("Example of processed vector (first 10 features):")
print(x_train_flat[0][:10])


Example of processed vector (first 10 features):
[0. 0. 0. 0. 0. 0. 0. 0. 0. 0.]

# -----
# 5 Unsupervised Learning: KMeans Clustering
# -----
# Cluster MNIST images into 10 clusters (digits 0-9)
kmeans = KMeans(n_clusters=10, random_state=42)
kmeans.fit(x_train_flat)


KMeans(n_clusters=10, random_state=42)
```

```
# Cluster assignments
clusters = kmeans.labels_


# -----
# 5a Show Cluster Centroids as Images
# -----
for i in range(10):
    centroid = kmeans.cluster_centers_[i].reshape(28,28)
    plt.imshow(centroid, cmap='gray')
    plt.title(f"Cluster {i} Centroid")
    plt.show()
```




Cluster 1 Centroid




Cluster 2 Centroid




Cluster 3 Centroid




Cluster 4 Centroid




Cluster 5 Centroid




Cluster 6 Centroid




Cluster 7 Centroid



Cluster 8 Centroid





```
# -----  
# 5b Map Clusters to Actual Digits (Insight)  
# -----  
# Count which label appears most in each cluster  
cluster_label_mapping = {}  
for i in range(10):  
    indices = np.where(clusters == i)[0]  
    most_common_label = np.bincount(y_train[indices]).argmax()  
    cluster_label_mapping[i] = most_common_label  
    print(f"Cluster {i} mostly contains digit: {most_common_label}")
```

```
Cluster 0 mostly contains digit: 4  
Cluster 1 mostly contains digit: 1  
Cluster 2 mostly contains digit: 7  
Cluster 3 mostly contains digit: 3  
Cluster 4 mostly contains digit: 9  
Cluster 5 mostly contains digit: 0  
Cluster 6 mostly contains digit: 1  
Cluster 7 mostly contains digit: 8  
Cluster 8 mostly contains digit: 6  
Cluster 9 mostly contains digit: 2
```

```
# -----  
# 6 Supervised Learning: Logistic Regression
```

```

# -----
# Logistic Regression classifier on flattened features
lr = LogisticRegression(max_iter=1000, solver='lbfgs',
multi_class='auto')
lr.fit(x_train_flat, y_train)

/usr/local/lib/python3.12/dist-packages/sklearn/linear_model/_logistic.py:1247: FutureWarning: 'multi_class' was deprecated in
version 1.5 and will be removed in 1.7. From then on, it will always
use 'multinomial'. Leave it to its default value to avoid this
warning.
    warnings.warn(

LogisticRegression(max_iter=1000, multi_class='auto')

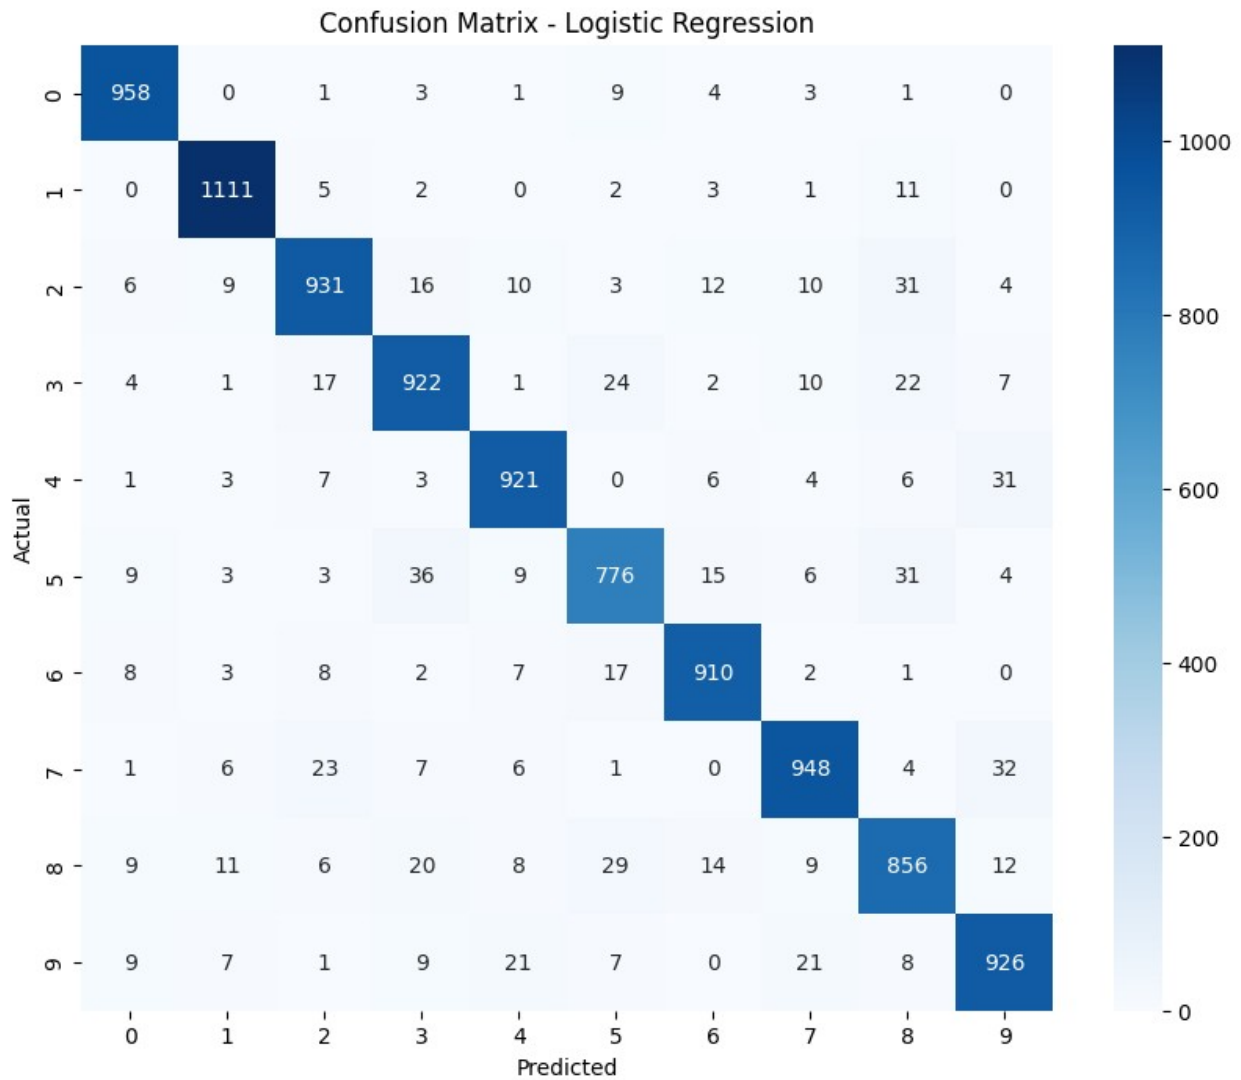
# Predict on test set
y_pred = lr.predict(x_test_flat)

# -----
# 6a Evaluate Model
# -----
accuracy = accuracy_score(y_test, y_pred)
print(f"Logistic Regression Accuracy: {accuracy*100:.2f}%")

Logistic Regression Accuracy: 92.59%


# Confusion matrix
cm = confusion_matrix(y_test, y_pred)
plt.figure(figsize=(10,8))
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues')
plt.title("Confusion Matrix - Logistic Regression")
plt.xlabel("Predicted")
plt.ylabel("Actual")
plt.show()

```




```
# -----
# 7 Some Predictions vs Actual
# -----
for i in range(5):
    plt.imshow(x_test[i], cmap='gray')
    plt.title(f"Actual: {y_test[i]}, Predicted: {y_pred[i]}")
    plt.show()
```


Actual: 7, Predicted: 7




Actual: 2, Predicted: 2



Actual: 1, Predicted: 1



Actual: 0, Predicted: 0



Actual: 4, Predicted: 4

